

Introduction: the Machine Learning Landscape

Xudong Liu

Assistant Professor

School of Computing

University of North Florida

Aug. 19, 2019

What is ML?

“A field of study that gives computers the ability to learn without being explicitly programmed.” – Arthur Samuel, 1959

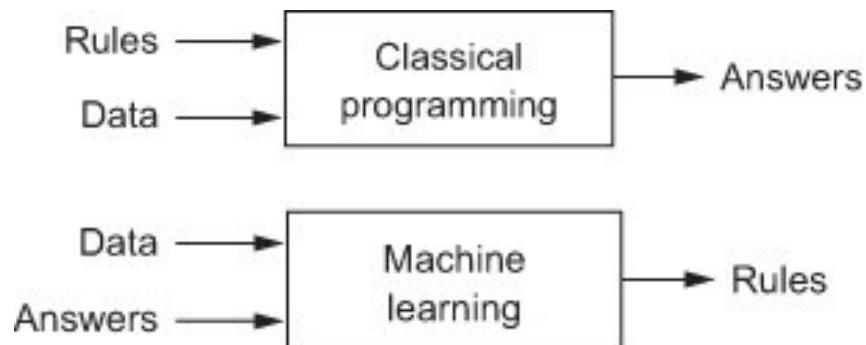
“A computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.” – Tom Mitchell, 1997

ML is a prominent sub-field in AI, “the new electricity.” – Andrew Ng

(<https://www.youtube.com/watch?v=21EiKfQYZXc>)

What is ML?

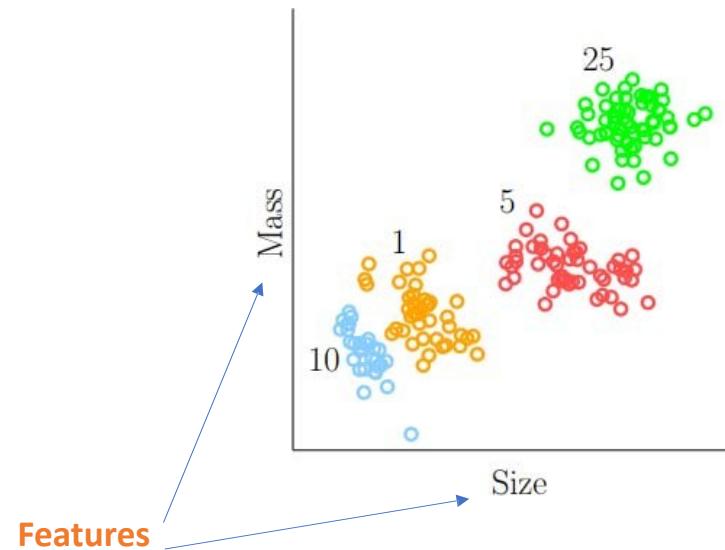
“A field of study that gives computers the ability to learn without being explicitly programmed.”



A machine-learning system is *trained* rather than explicitly programmed.

Types of ML Systems

Supervised Learning - Training Data contains desired solutions, or *labels*



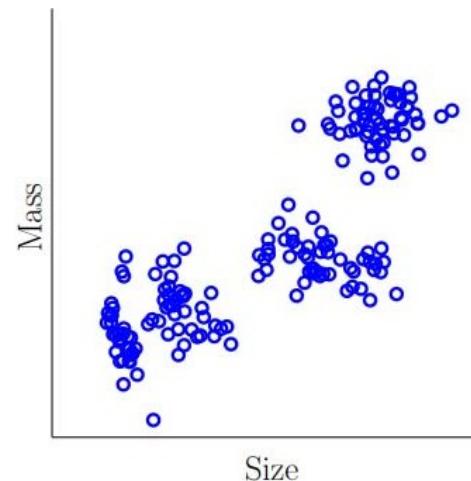
Types of ML Systems

Supervised Learning - Training Data contains desired solutions, or *labels*

1. K-Nearest Neighbors
2. Linear Regression
3. Logistic Regression
4. Support Vector Machines
5. Decision Trees and Random Forests
6. Neural Networks

Types of ML Systems

Unsupervised Learning - Training Data is *unlabeled*



Types of ML Systems

Unsupervised Learning - Training Data is *unlabeled*

1. K-Means (Clustering)
2. Principal Component Analysis (Dimensionality Reduction)
3. Deep Neural Networks

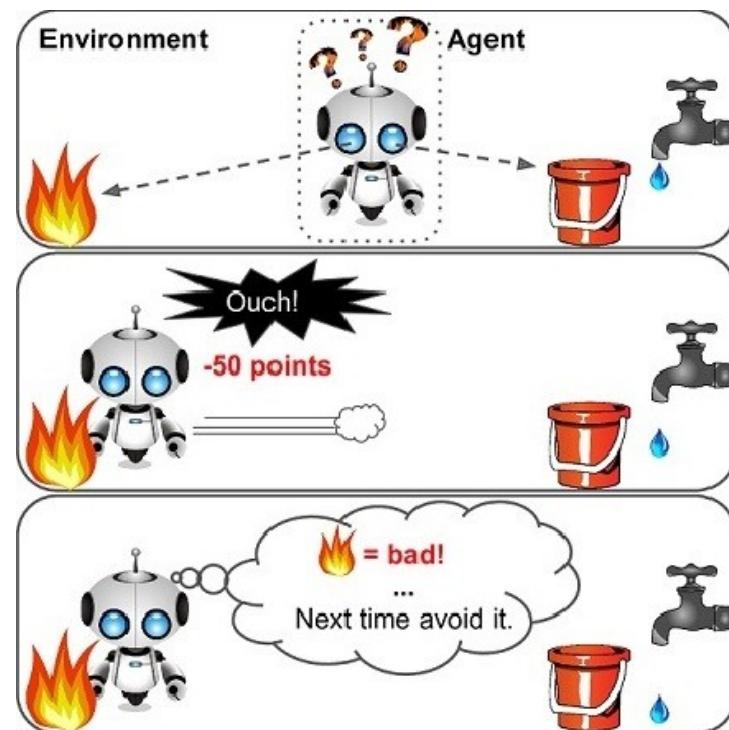
Types of ML Systems

Reinforcement Learning - Training Data does not contain target output, but instead contains **some** possible output together with a measure of how **good** that output is.

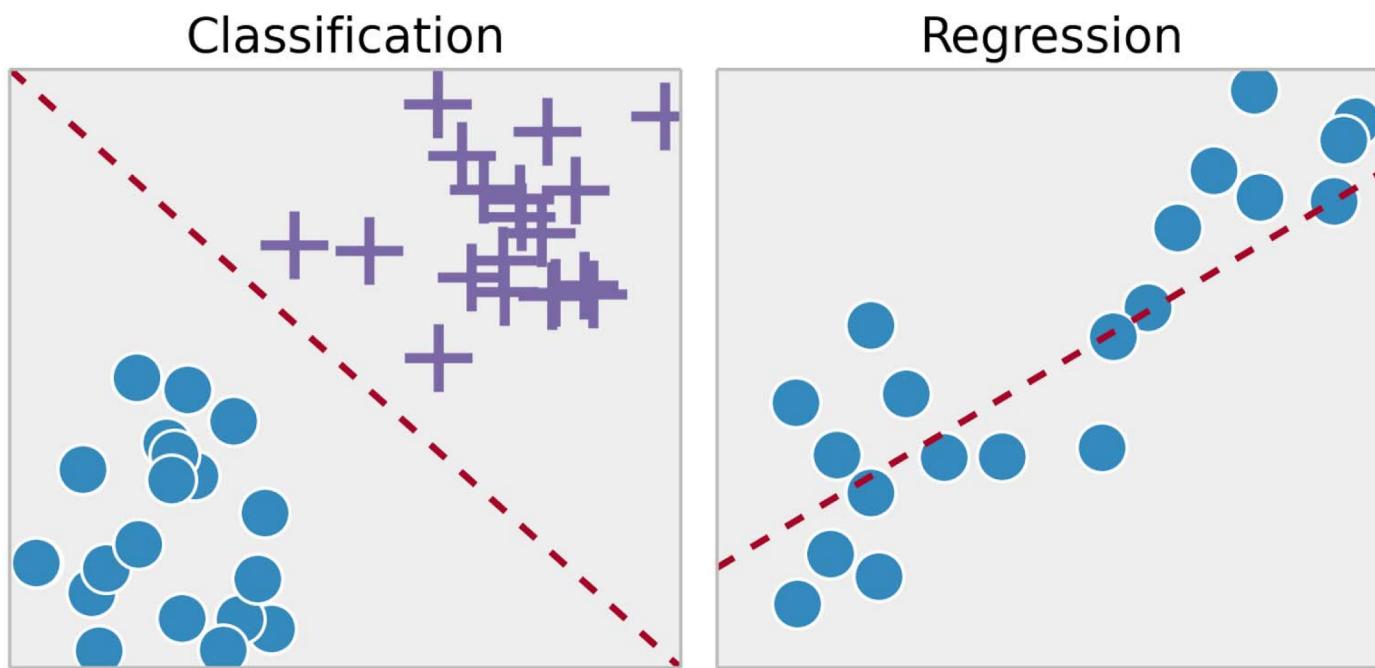
<input>, <correct output>



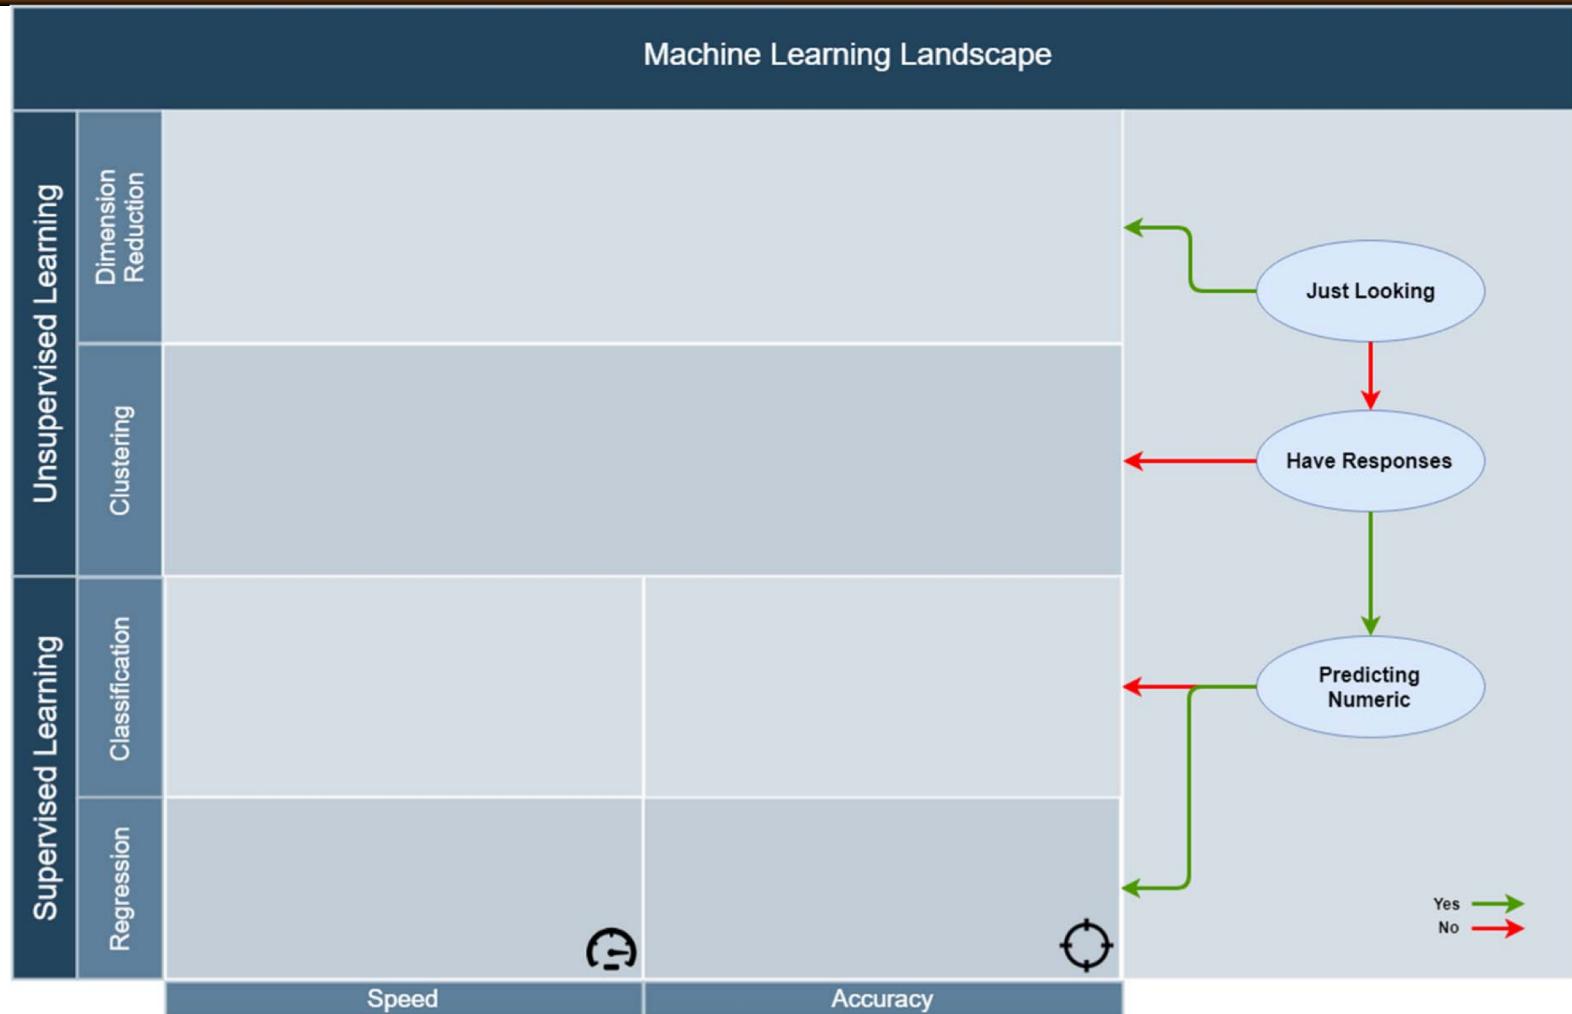
<input>, <some output>, <grade for this output>



Classification vs Regression



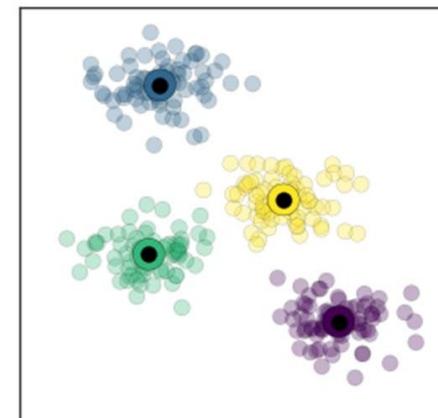
ML Landscape



Unsupervised Learning - Clustering

Clustering

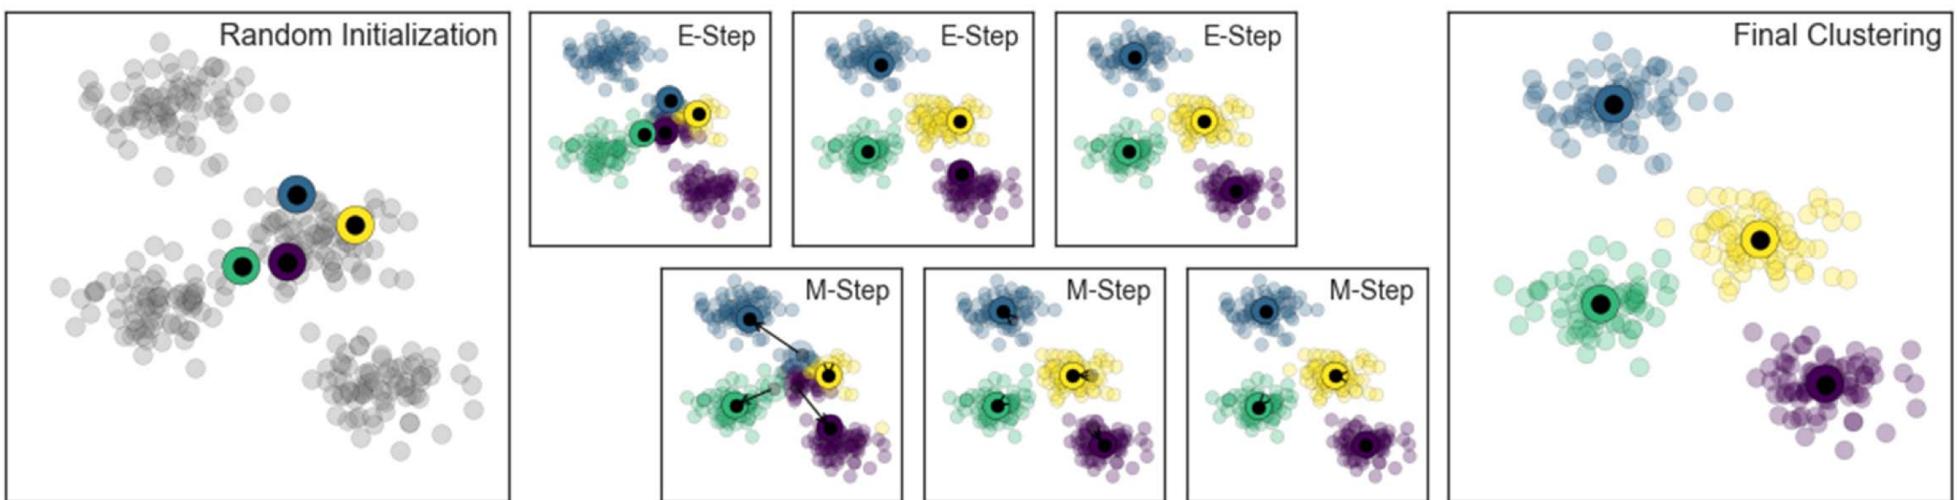
- Color “clusters” of points in a homogenous cloud of data.
- **Use Cases**
 - Behavioral Segmentation in Marketing
 - Useful as a pre-processing step before applying other classification algorithms.
 - Cluster ID could be added as feature for each data point.



Unsupervised Learning - Clustering

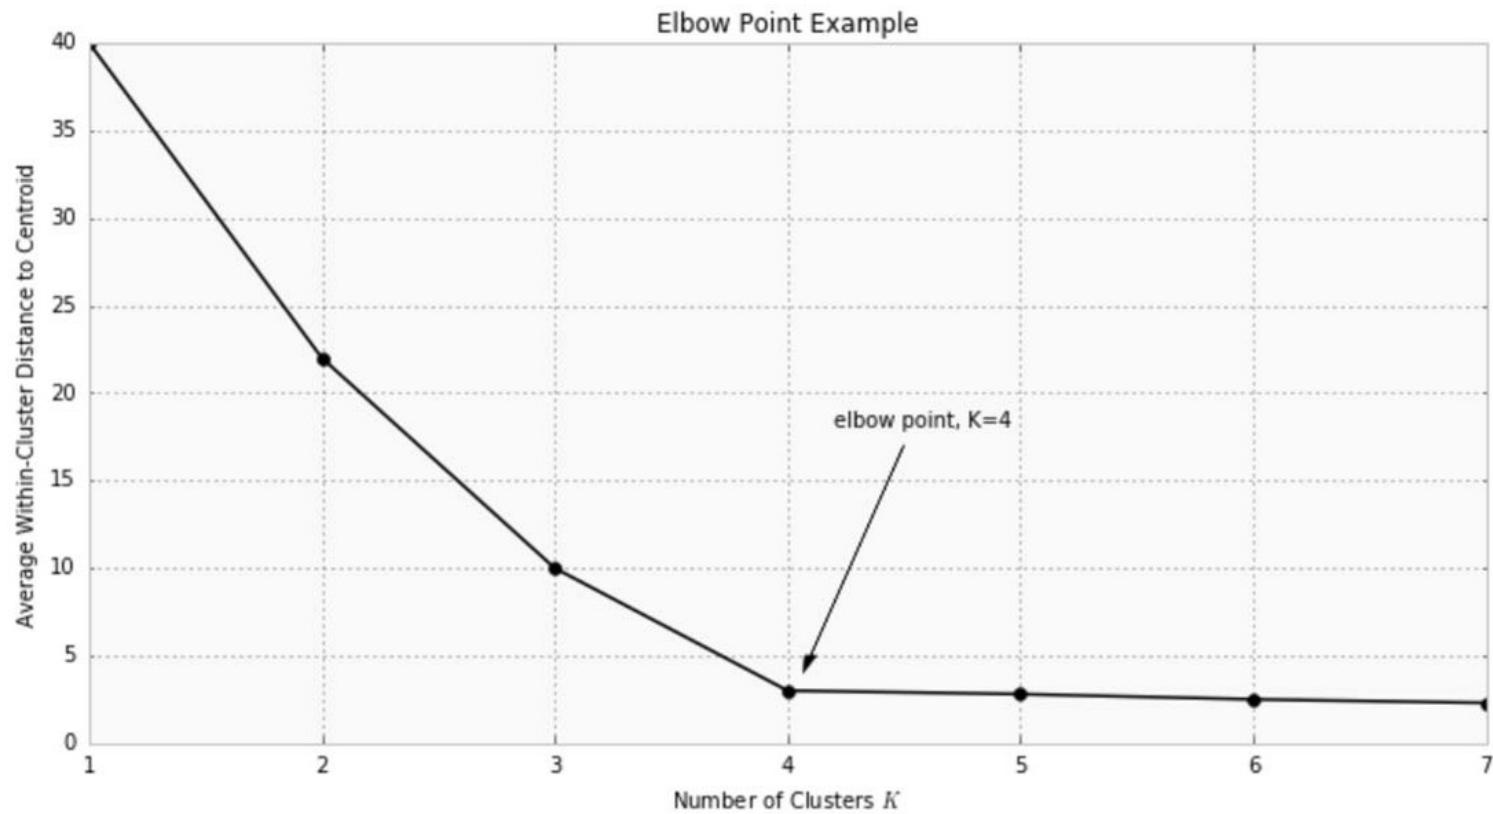
k-means Algorithm

- Guess some cluster centers
- Repeat until converged
 - *E-Step*: assign points to the nearest cluster center
 - *M-Step*: set the cluster centers to the mean

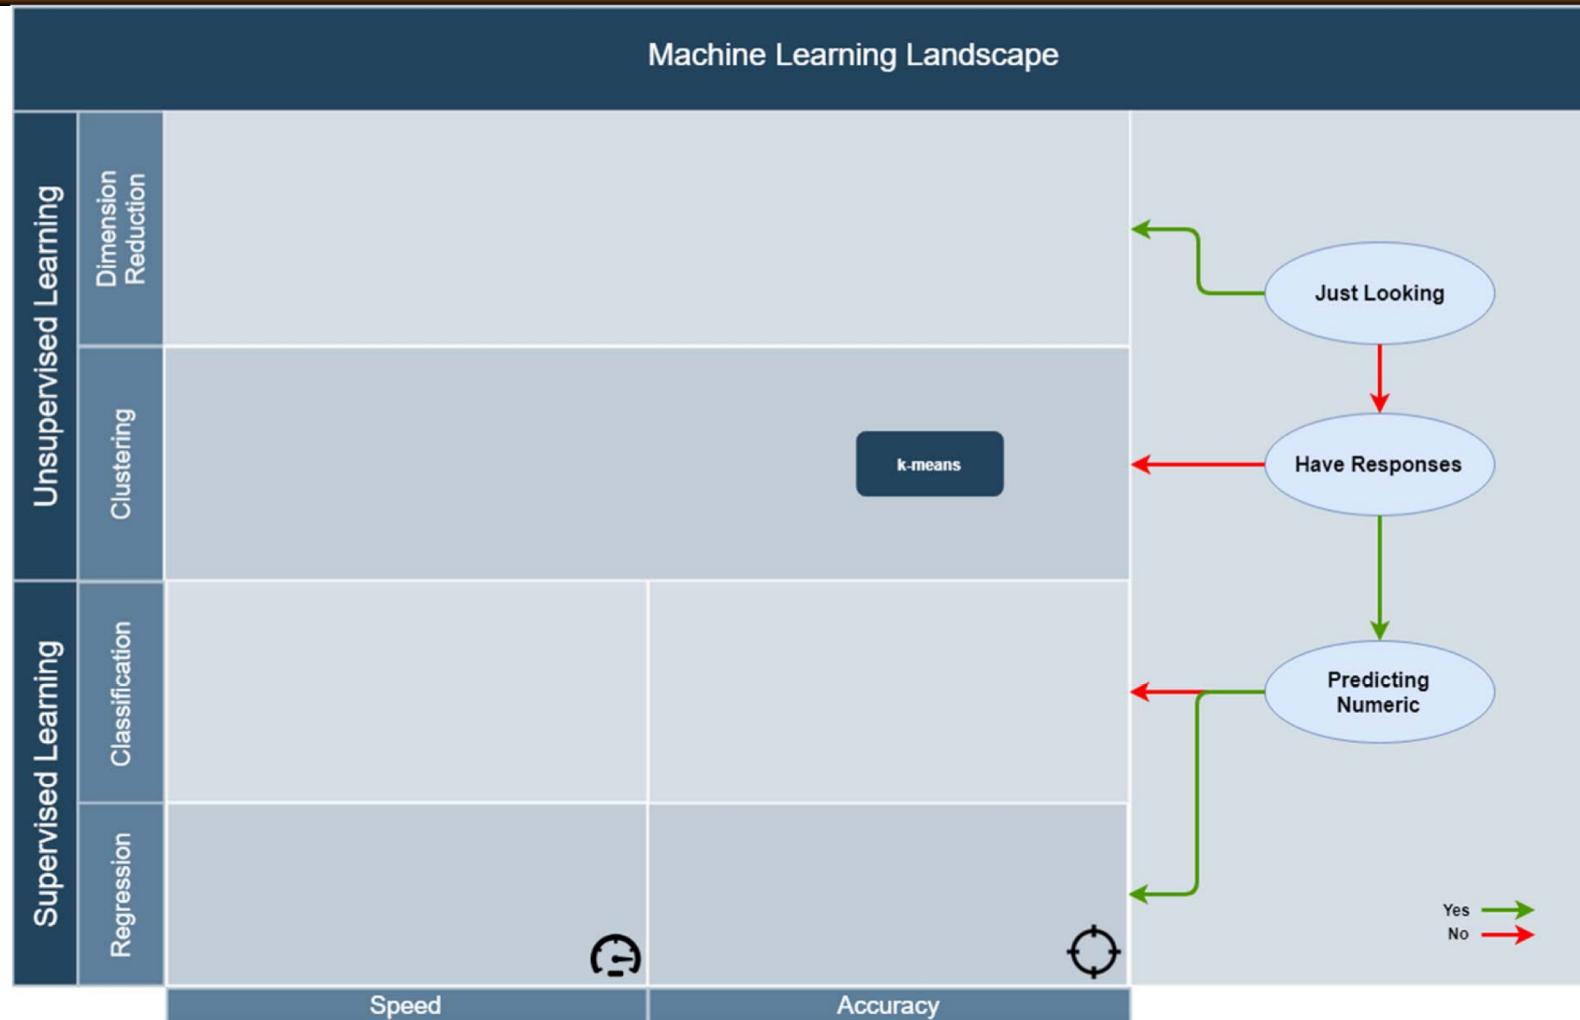


Unsupervised Learning - Clustering

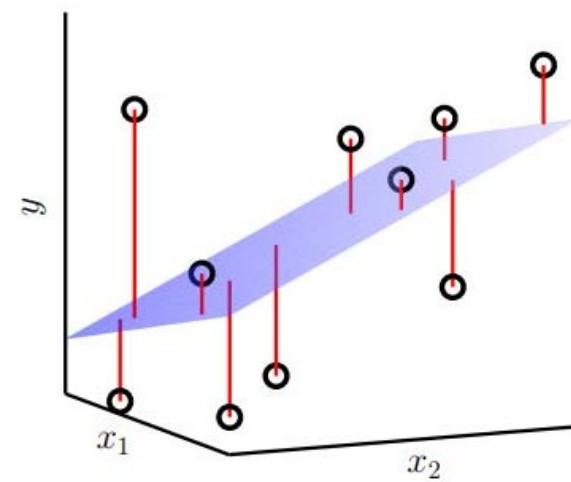
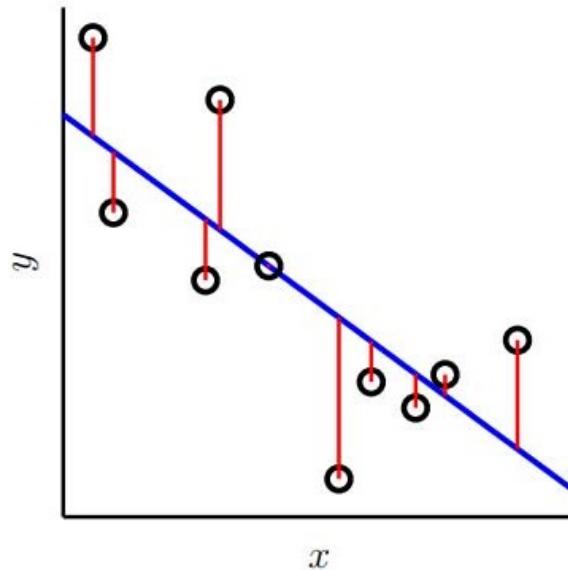
Choosing k



ML Landscape



Linear Regression



Linear Regression

$$X = \begin{bmatrix} 1 & x_1^1 & x_2^1 & \dots & x_n^1 \\ 1 & x_1^2 & x_2^2 & \dots & x_n^2 \\ \vdots & \vdots & \dots & \dots & \vdots \\ 1 & x_1^m & x_2^m & \dots & x_n^m \end{bmatrix}$$

$$y = \begin{bmatrix} y^1 \\ y^2 \\ \dots \\ y^m \end{bmatrix}$$

$$\Theta = \begin{bmatrix} \Theta_0 \\ \Theta_1 \\ \Theta_2 \\ \vdots \\ \Theta_n \end{bmatrix} \quad \hat{y} = \begin{bmatrix} \hat{y}^1 \\ \hat{y}^2 \\ \dots \\ \hat{y}^m \end{bmatrix}$$

Define a **Hypothesis**

$$h_{\Theta}(X) = \hat{y} = X \Theta$$

Define a **Cost Function** (a measure of how bad we're doing)

$$MSE(X, h_{\Theta}) = \frac{1}{n} \sum_{i=1}^n [y^i - \hat{y}^i]^2$$

Repeat until convergence:

- Calculate Cost Function on chosen Θ
- Calculate slope of Cost Function
- Tweak Θ so as to move downhill (reduce Cost Function value)

Θ is now optimized for our training data.

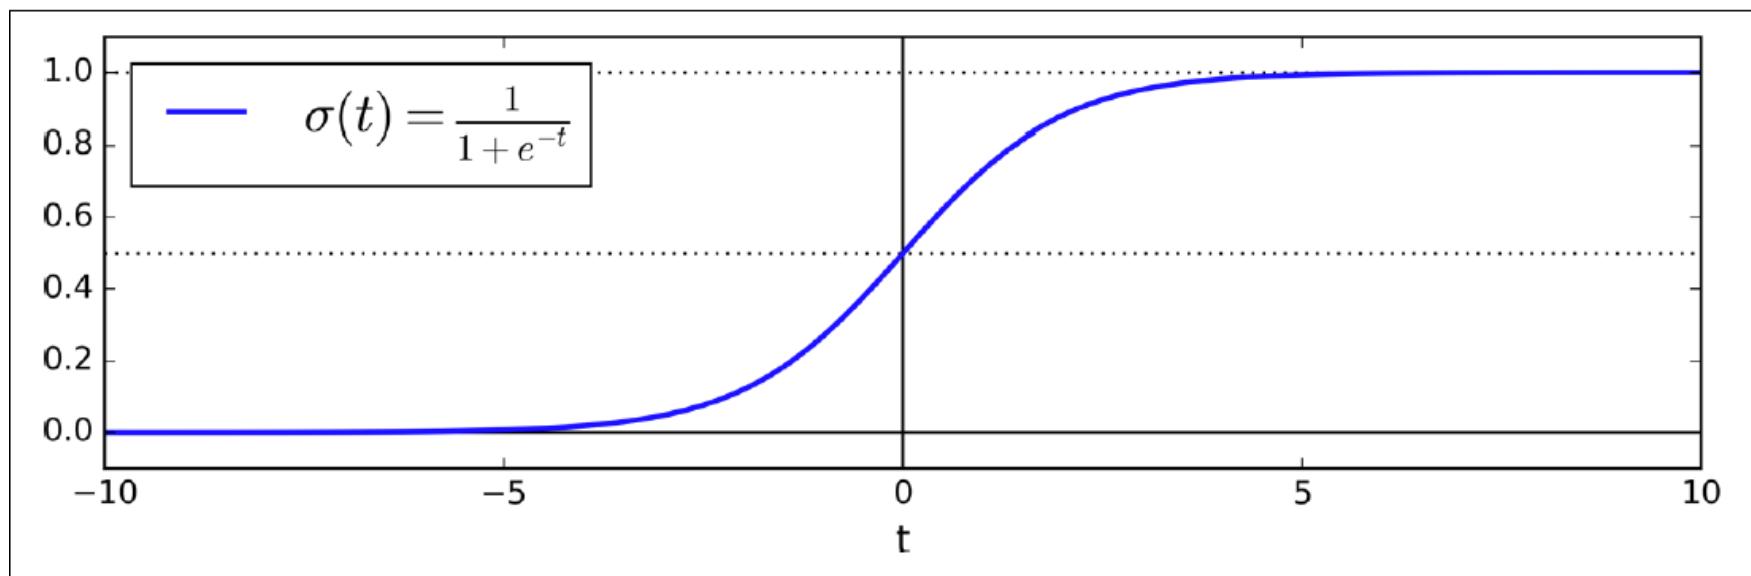
Logistic Regression

Used to estimate the probability that an instance belongs to a particular class.

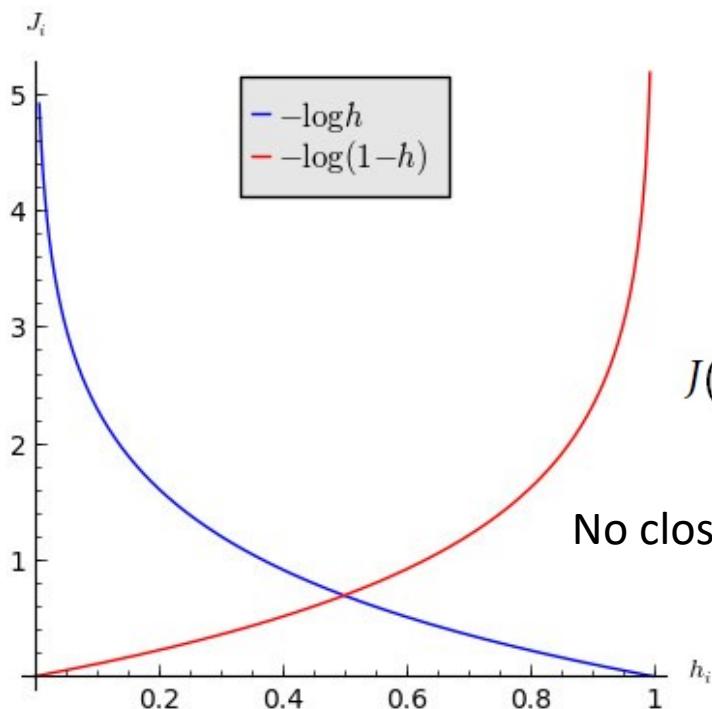
$$\hat{p} = h_{\theta}(\mathbf{x}) = \sigma(\boldsymbol{\theta}^T \cdot \mathbf{x})$$

Logistic Regression

$$\sigma(t) = \frac{1}{1 + \exp(-t)}$$



Logistic Regression

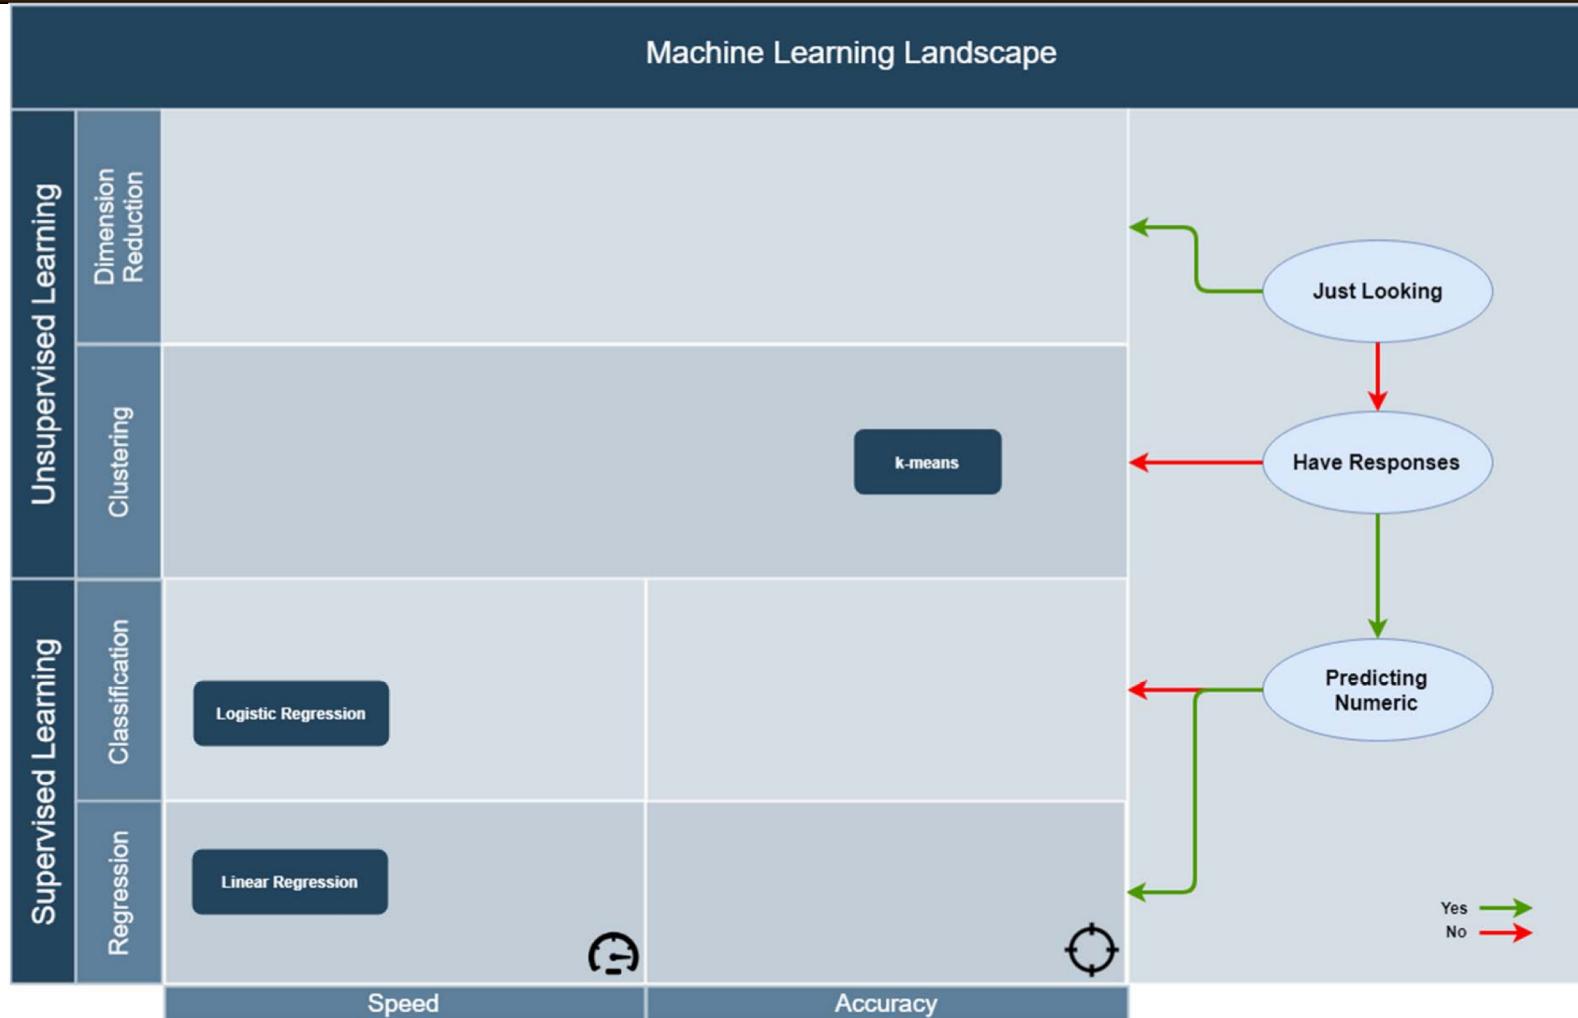


$$c(\theta) = \begin{cases} -\log(\hat{p}) & \text{if } y = 1, \\ -\log(1 - \hat{p}) & \text{if } y = 0. \end{cases}$$

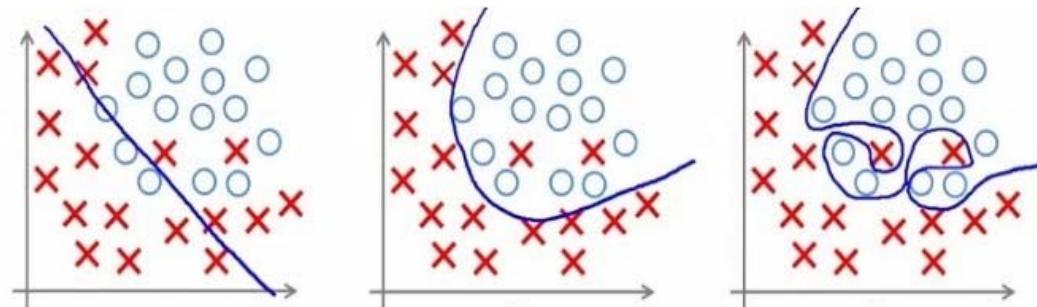
$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(\hat{p}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{p}^{(i)}) \right]$$

No closed-form solution, but we can use **Gradient Descent!**

ML Landscape



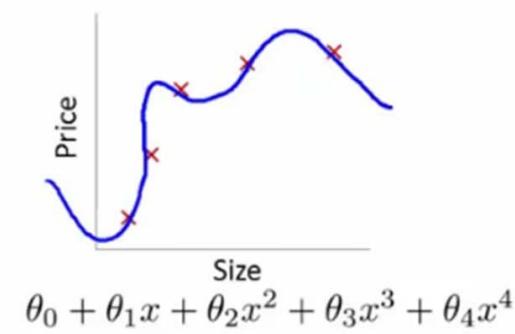
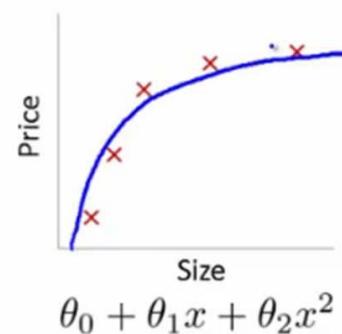
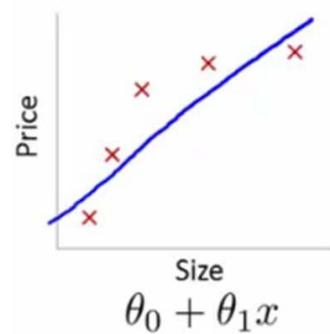
Overfitting and Underfitting



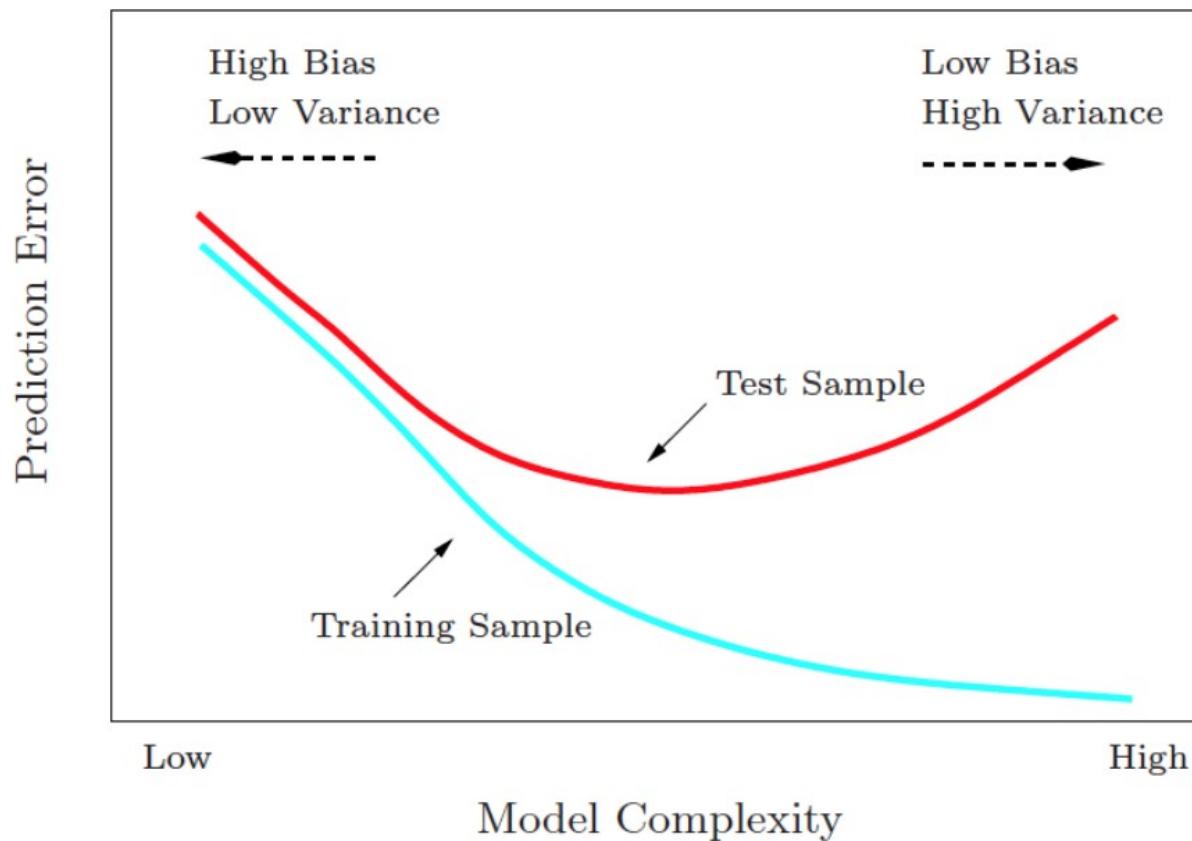
Under-fitting

Appropriate-fitting

Over-fitting



Bias-Variance Tradeoff



Regularization

How to ensure that we're not *overfitting* to our training data?

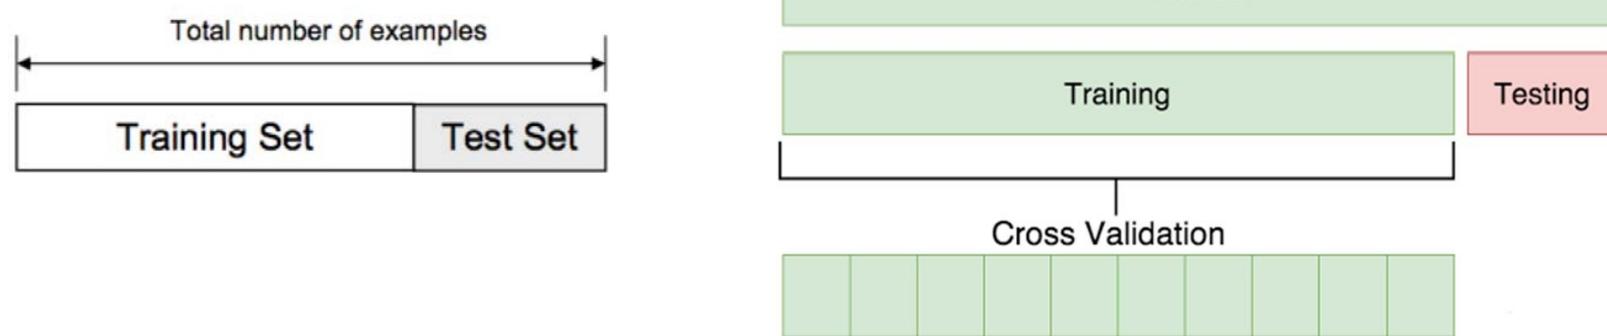
Impose a small penalty on model complexity.

|1 penalty (Lasso Regression)

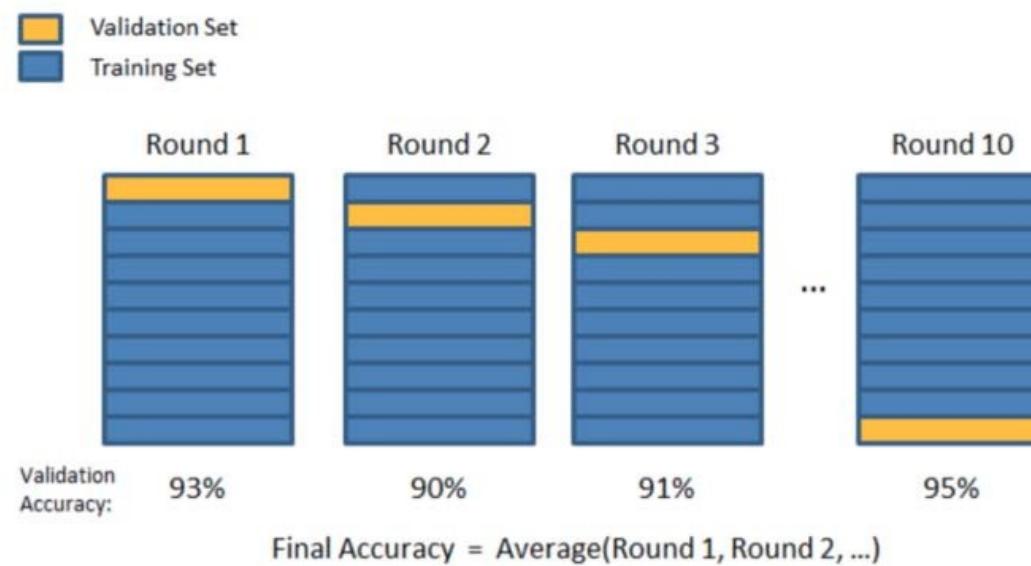
|2 penalty (Ridge Regression)

$$J(\theta) = \text{MSE}(\theta) + \alpha \sum_{i=1}^n |\theta_i| \quad J(\theta) = \text{MSE}(\theta) + \alpha \frac{1}{2} \sum_{i=1}^n \theta_i^2$$

Testing and Validation



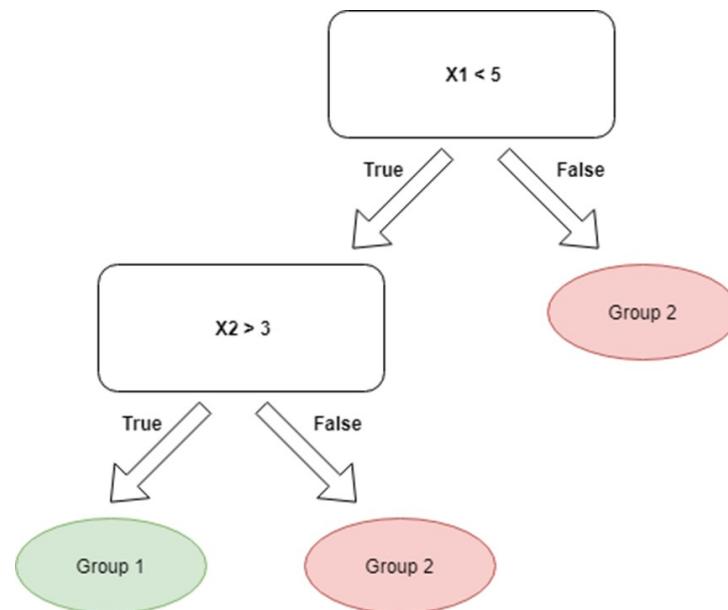
K-fold Cross Validation



Decision Tree

Basic Idea

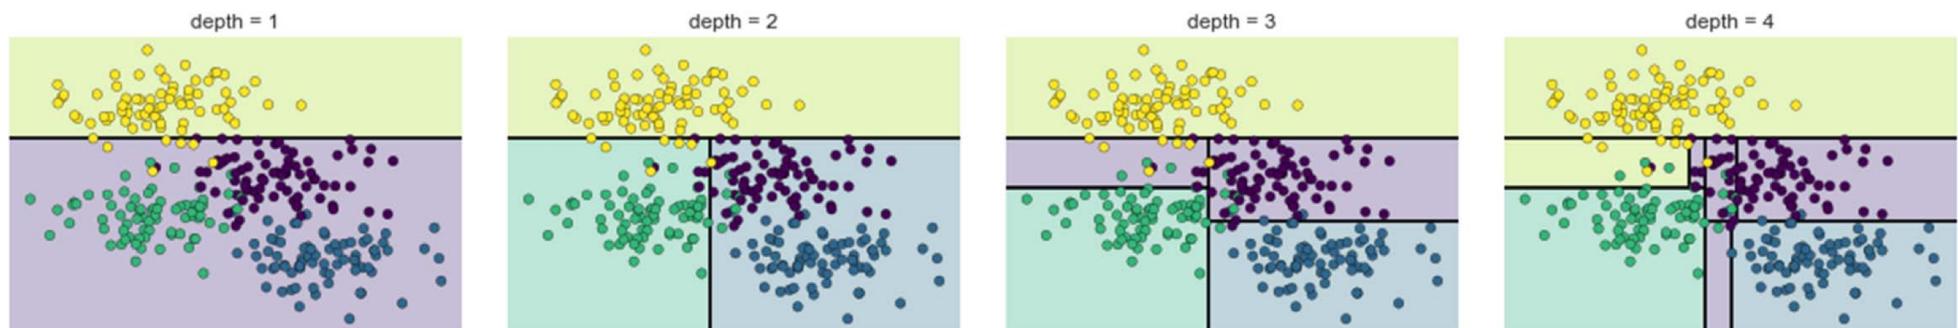
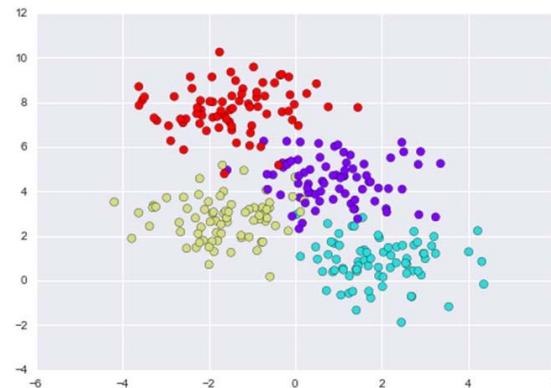
- Construct a tree to ask a series of questions from your data.



Decision Tree

Let's see how it works on a *real* dataset.

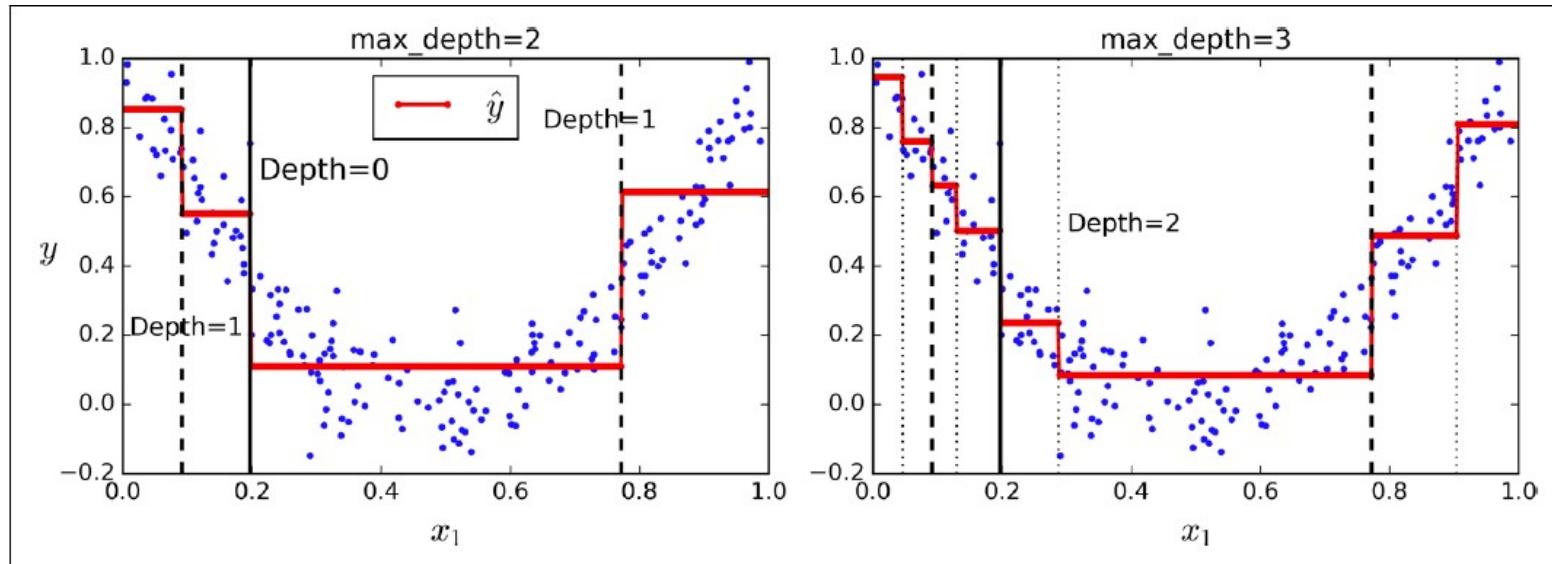
- Boundaries always are **axis-parallel**



Decision Tree

Decision Trees can be used for regression!

- Minimize MSE instead of impurity.



Decision Tree

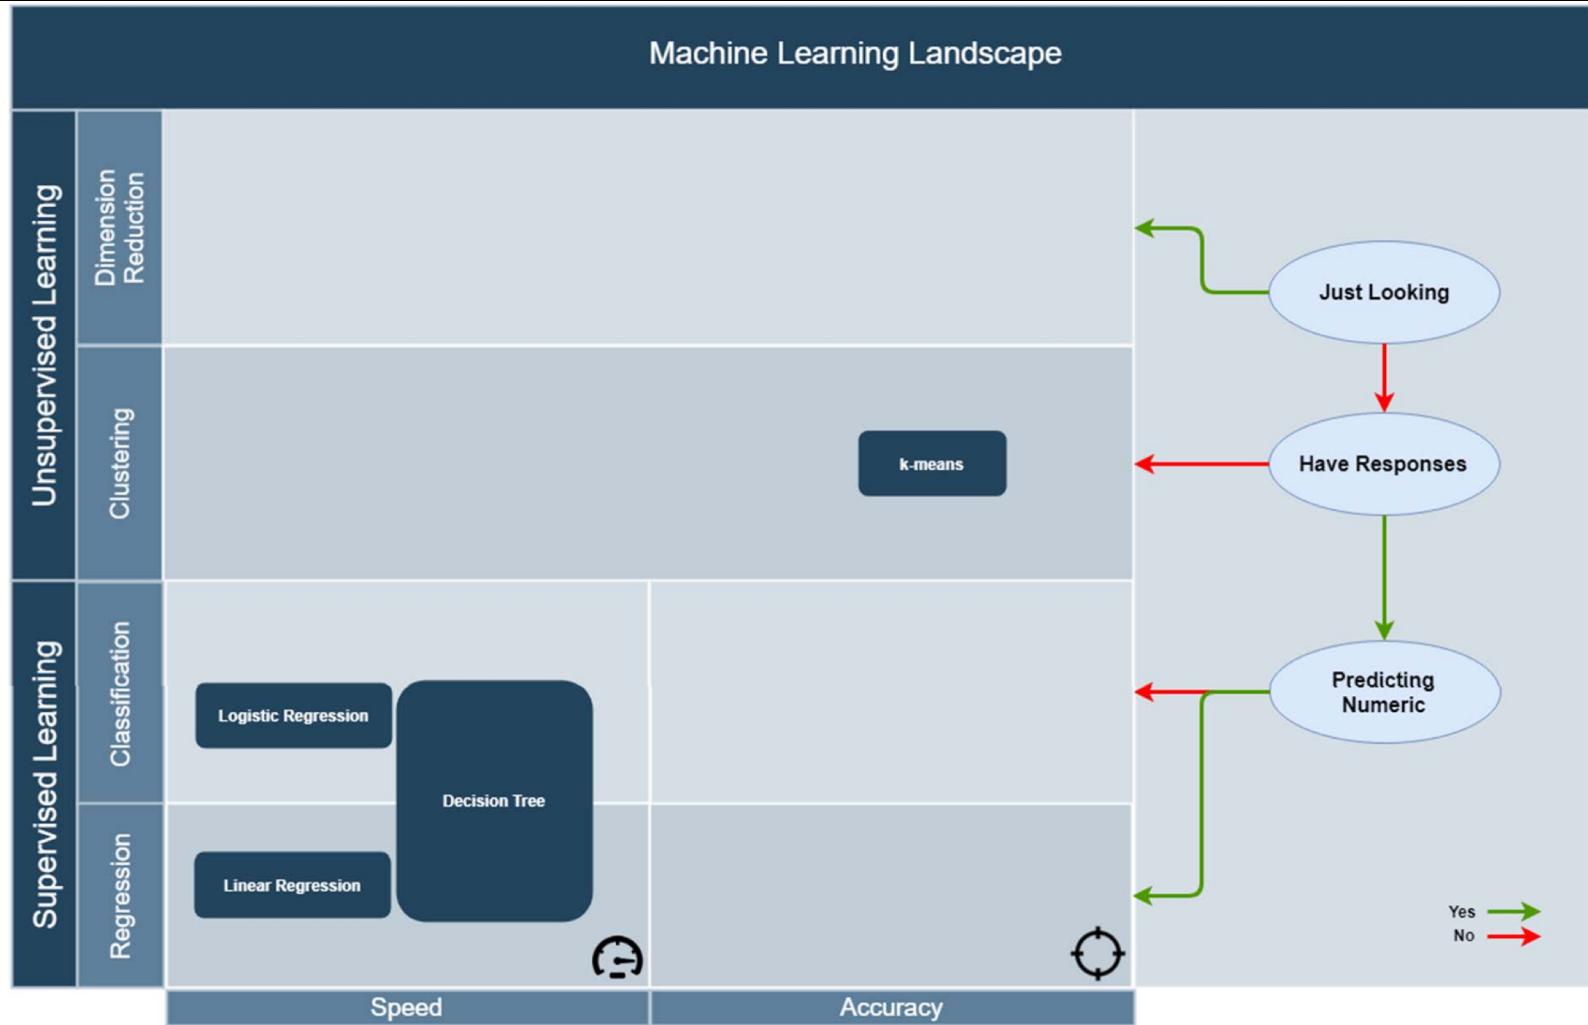
Advantages

- *White Box* – easily interpretable

Disadvantages

- Prone to overfitting
 - Regularize by setting maximum depth
- Comes up only with orthogonal boundaries
 - Sensitive to training set rotation – Use PCA!

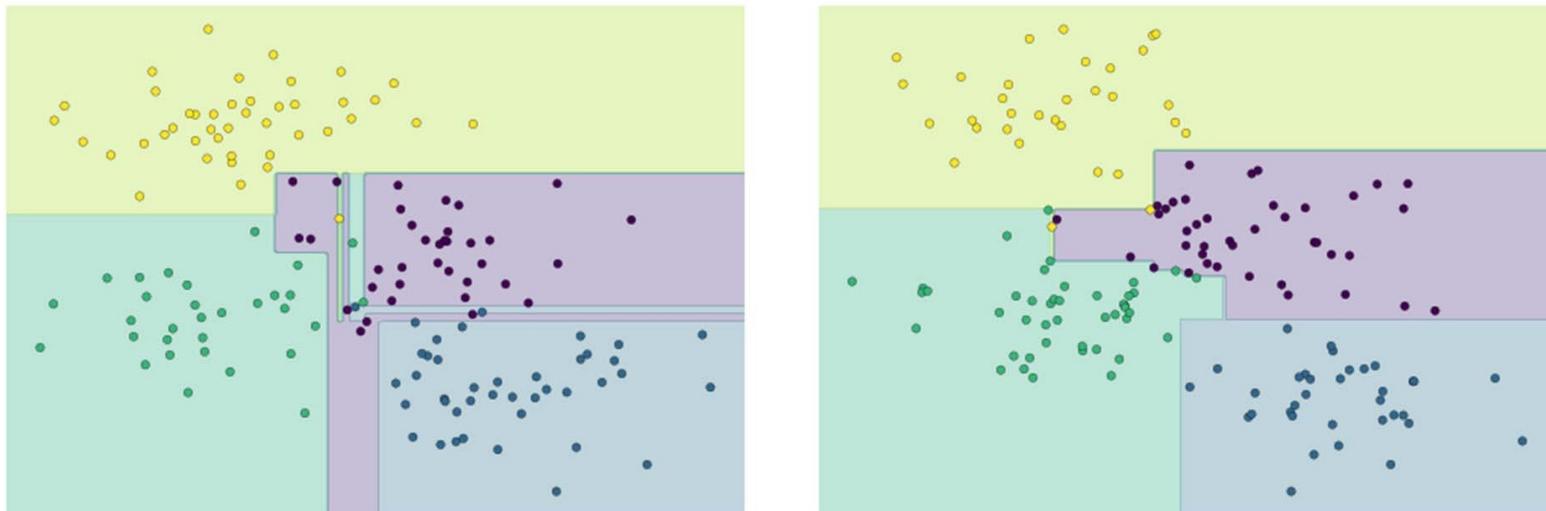
ML Landscape



Ensemble Methods

Basic Idea

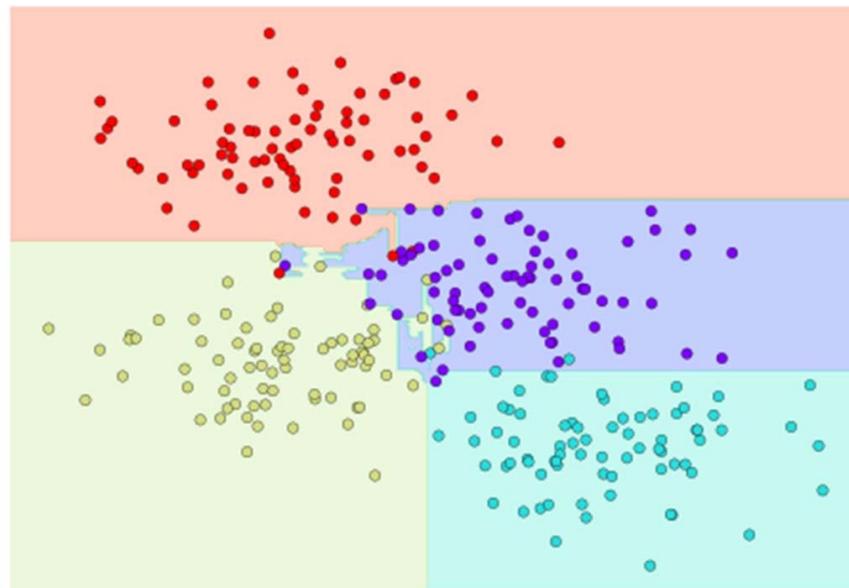
- Two Decision Trees by themselves may overfit. But combining their predictions may be a good idea!



Ensemble Methods

Basic Idea

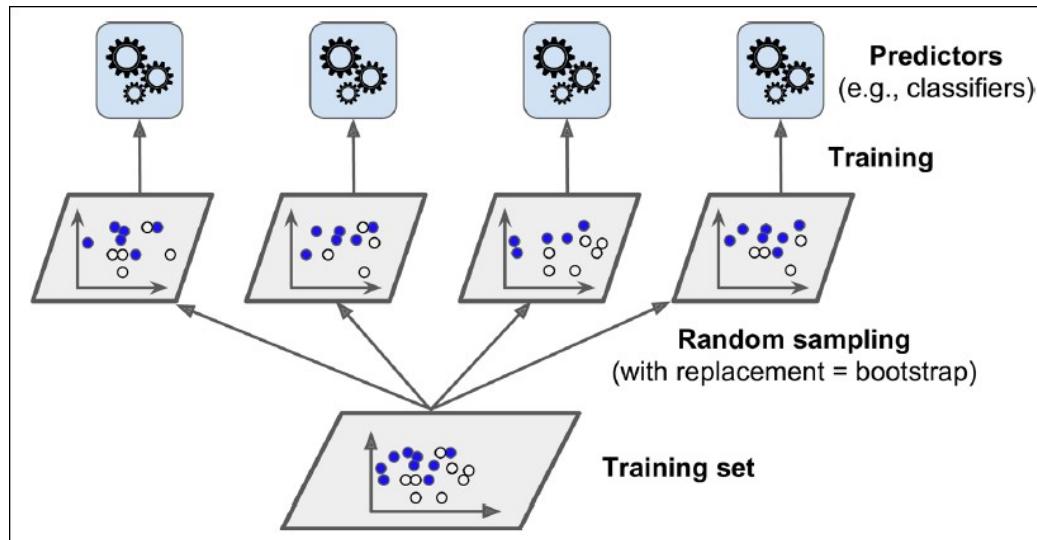
- Two Decision Trees by themselves may overfit. But combining their predictions may be a good idea!



Bagging

Bagging = Bootstrap Aggregation

- Use the same training algorithm for every predictor, but train them on different random subsets of the training set.

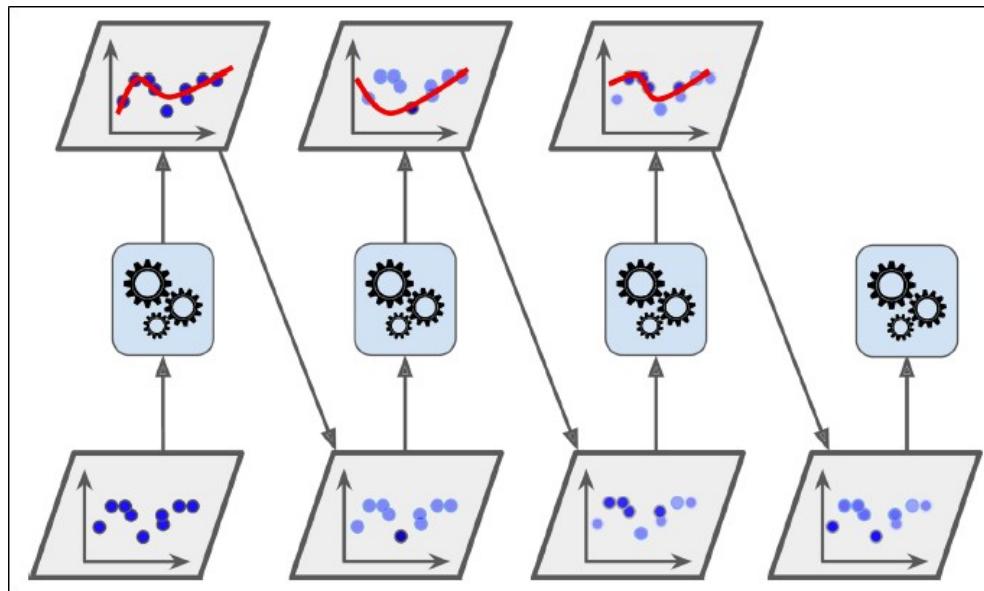


Random Forest is an *Ensemble* of Decision Trees, generally trained via the *bagging* method.

Boosting

Basic Idea

- Train several weak learners *sequentially*, each trying to correct the errors made by its predecessor.



Adaptive Boosting (ADABoost)

- Give more relative weight to the misclassified instances.

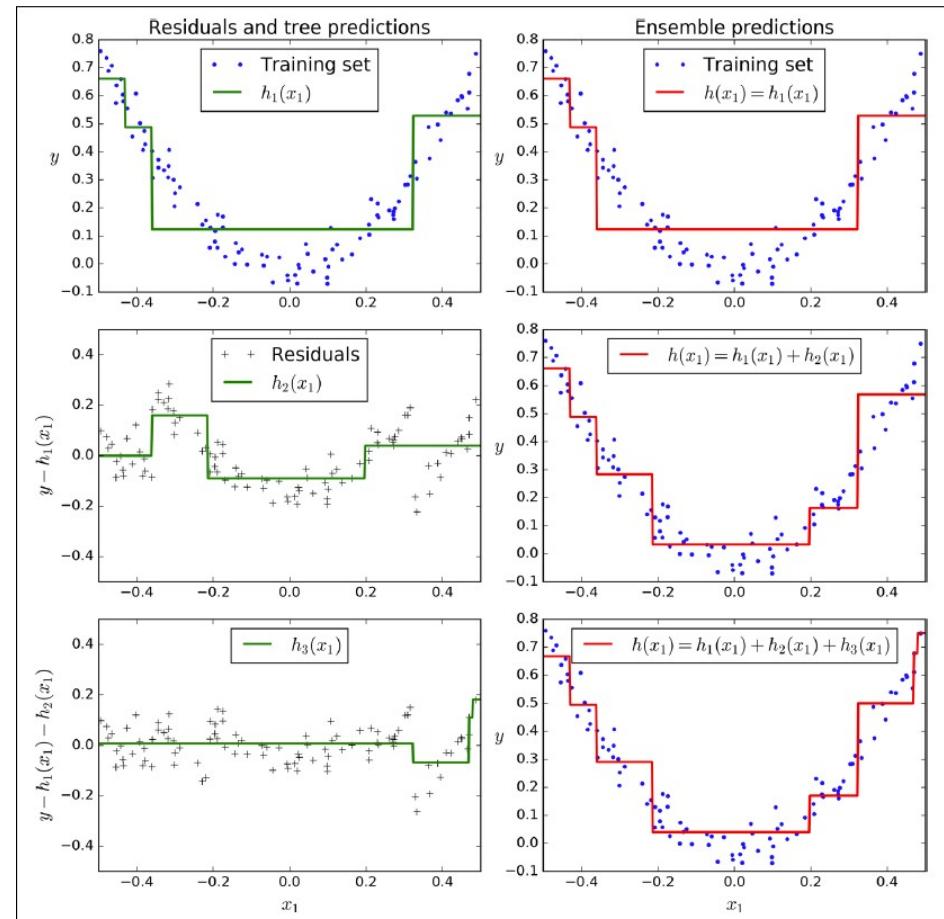
Boosting

Gradient Boosting

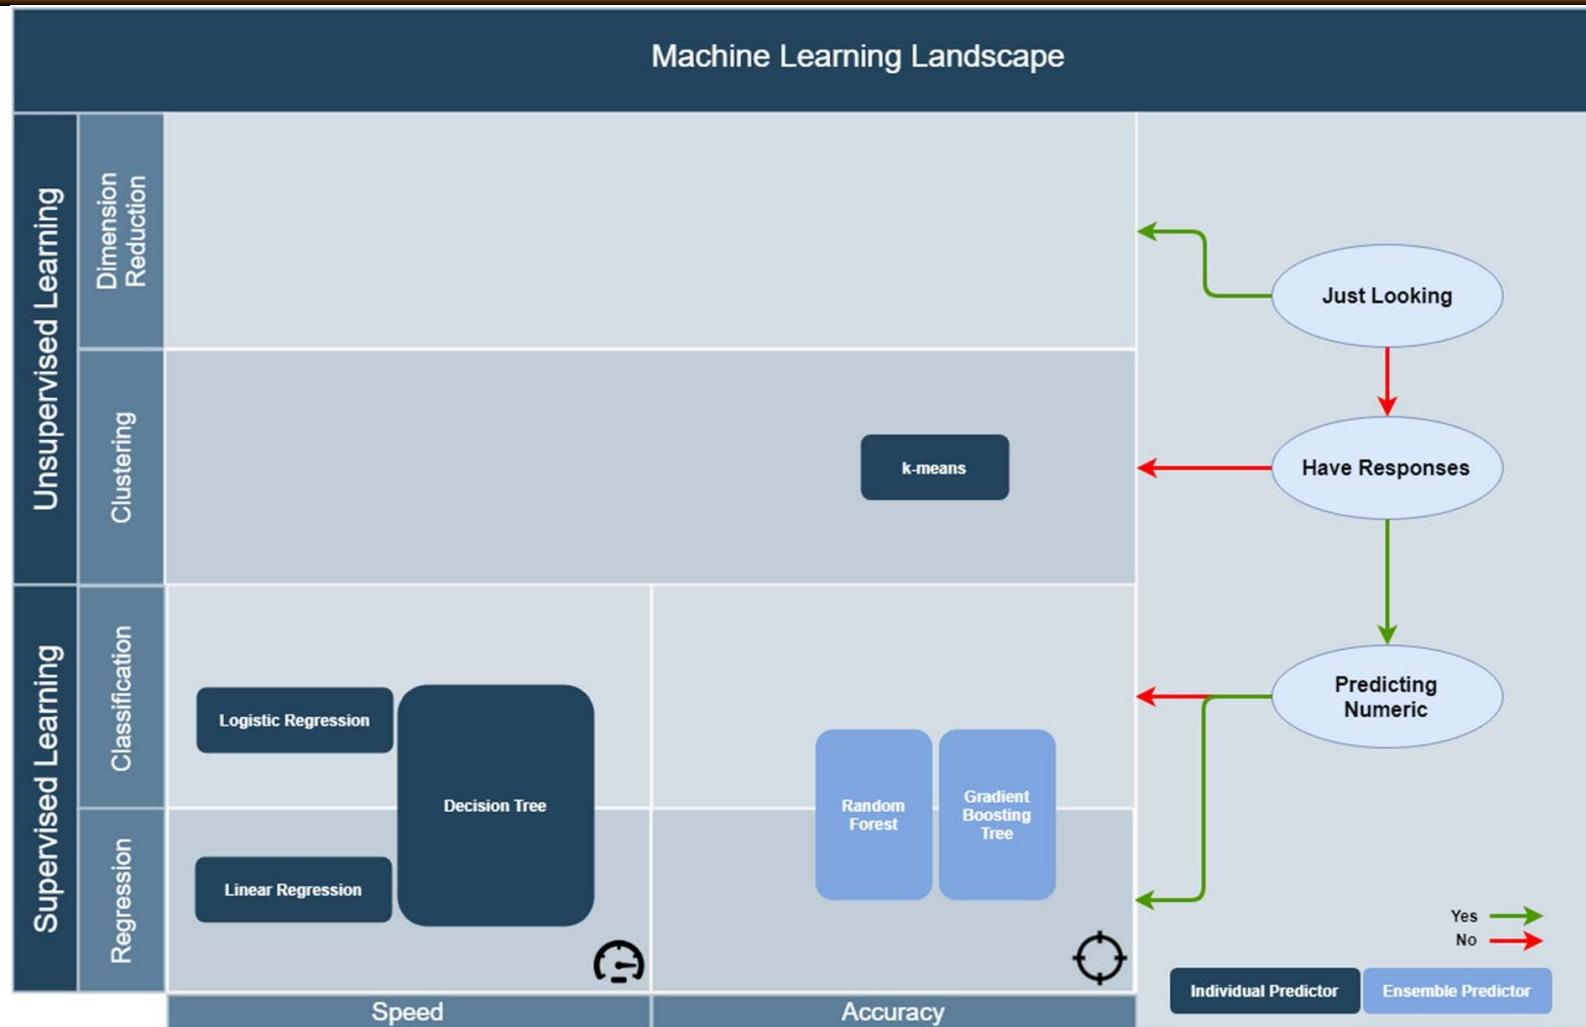
- Try to fit a new predictor to the *residual errors* made by the previous predictor.

Best Performance

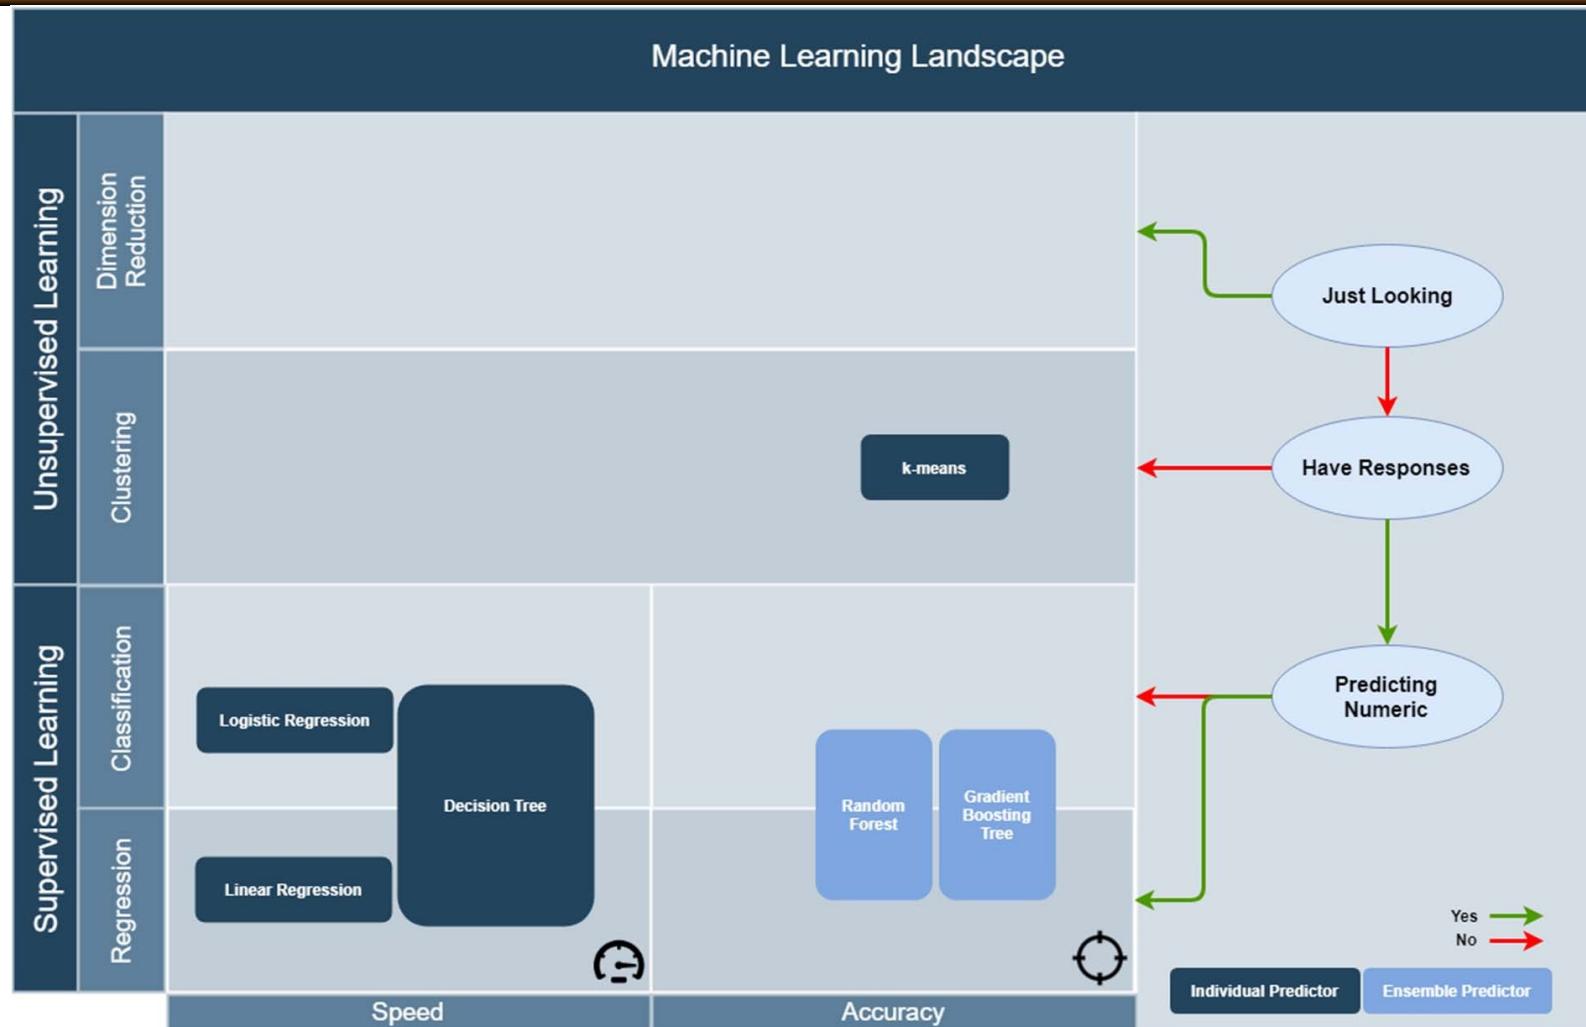
- Random Forests and Gradient Boosting Methods (implemented in the *xgBoost* library) have been winning most competitions on Kaggle recently on structured data.
- Deep Learning (especially Convolutional Networks) is the clear winner for unstructured data problems (perception/speech/vision etc.)



ML Landscape

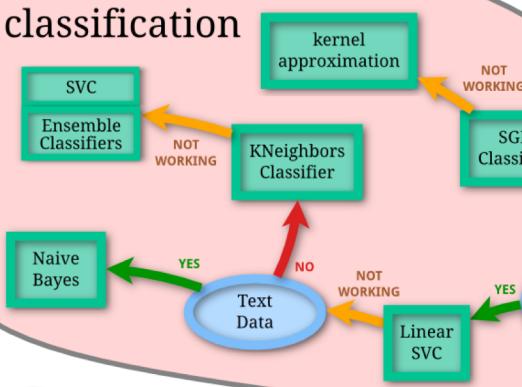


ML Landscape

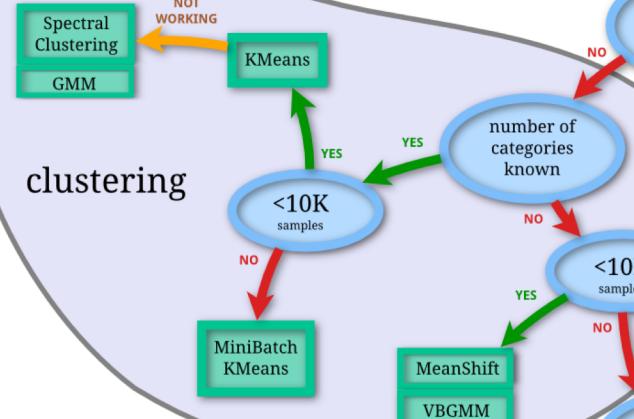


ML Landscape

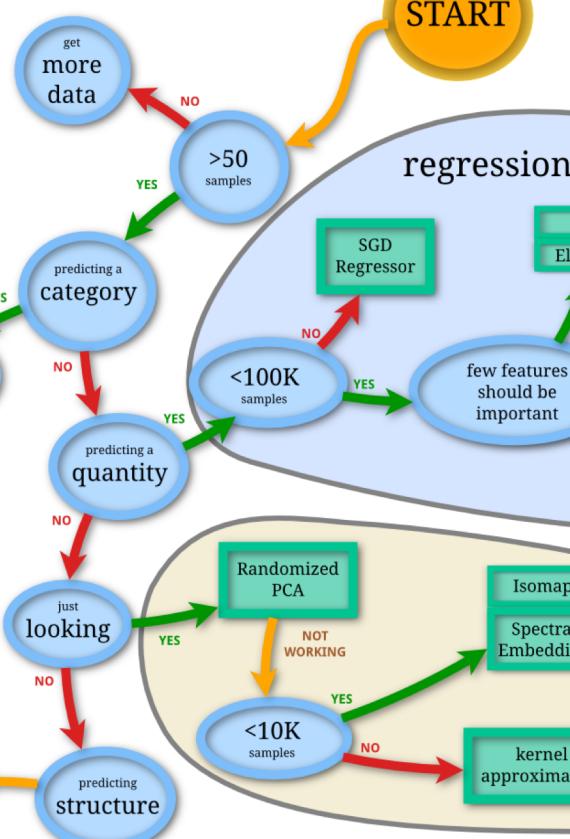
classification



clustering



scikit-learn
algorithm cheat-sheet



dimensionality
reduction

Back

scikit
learn